

COMP 737011 - Memory Safety and Programming Language Design

Lecture 12: Dynamic Analysis of Rust Code

徐 辉

xuh@fudan.edu.cn



Outline

- 1. Background of Dynamic Analysis
- 2. Blackbox Fuzzing
- 3. Fuzz Target Generation
- 4. Dynamic Symbolic Execution

1. Background of Dynamic Analysis

Any Bugs in genvec()?

```
fn genvec(s:&mut String)->Vec<u8>{
    let mut v = Vec::new();
    let l = s.len();
    if l%10 == 0 {
        let ptr = s.as_mut_ptr();
        unsafe{
            v = Vec::from_raw_parts(ptr,l,s.len());
        }
    }
    return v;
}


fn main(){
    let args:Vec<String> = env::args().collect();
    let mut s = String::from(&args[1]);
    let v = genvec(&mut s);
    println!("{:?}",v);
}
```

Challenge of Bug Detection



- Testing?
 - Bugs are on uncommon execution paths.
 - How to achieve high coverage (code/path)
 - define test cases manually?
 - automated test generation
 - Executing all paths is almost infeasible
- How to test Library APIs?

Programs as Computation Trees

- A control-flow graph with loops unrolled.
- Each node is a a conditional statement (state)
- Each edge is a sequence of sequential statements
- Each path represents an equivalence class of inputs.




Control-flow graph



Computation tree

Example of Computation Tree

```
fn genvec(s:&mut String)->Vec<u8>{  
  let mut v = Vec::new();  
  let l = s.len();  
  if l%10 == 0 {  
    let ptr = s.as_mut_ptr();  
    unsafe{  
      v = Vec::from_raw_parts(ptr,l,s.len());  
    }  
  }  
  return v;  
}
```



How to Traverse the Tree?

- Random testing or fuzzing:
 - generate random inputs to execute the program.
 - probability of reaching error could be very small
- Symbolic execution:
 - Collect the constraints for each path
 - Use constraint solvers to find solutions

```
if x == 94389 {  
    unreachable!();  
}
```

Probability = $1/2^{32}$

2. Blackbox Fuzzing

Evolutionary Fuzzing with AFL

- AFL is the most famous fuzzing tool
 - <http://lcamtuf.coredump.cx/afl/>
 - Many followup tools available, e.g., afl++
- Use an genetic approach to learn interesting mutations:
 - trigger code code => keep it
 - cannot trigger new code => discard

american fuzzy lop 0.47b (readpng)			
process timing		overall results	
run time : 0 days, 0 hrs, 4 min, 43 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 26 sec		total paths : 195	
last uniq crash : none seen yet		uniq crashes : 0	
last uniq hang : 0 days, 0 hrs, 1 min, 51 sec		uniq hangs : 1	
cycle progress		map coverage	
now processing : 38 (19.49%)		map density : 1217 (7.43%)	
paths timed out : 0 (0.00%)		count coverage : 2.55 bits/tuple	
stage progress		findings in depth	
now trying : interest 32/8		favored paths : 128 (65.64%)	
stage execs : 0/9990 (0.00%)		new edges on : 85 (43.59%)	
total execs : 654k		total crashes : 0 (0 unique)	
exec speed : 2306/sec		total hangs : 1 (1 unique)	
fuzzing strategy yields		path geometry	
bit flips : 88/14.4k, 6/14.4k, 6/14.4k		levels : 3	
byte flips : 0/1804, 0/1786, 1/1750		pending : 178	
arithmetics : 31/126k, 3/45.6k, 1/17.8k		pend fav : 114	
known ints : 1/15.8k, 4/65.8k, 6/78.2k		imported : 0	
havoc : 34/254k, 0/0		variable : 0	
trim : 2876 B/931 (61.45% gain)		latent : 0	



Fuzz Rust with AFL

1. Install AFL

<https://rust-fuzz.github.io/book/introduction.html>

2. Create a new project

```
#: cargo install afl
```

```
#: cargo new fuzztarget
```

3. Add deps in Cargo.toml

```
[dependencies] afl = "*"
```

4. fuzz target

```
#[macro_use]
extern crate afl;
fn genvec(s:&mut String)->Vec<u8>{ ...}
fn main(){
    fuzz!(|data: &[u8]| {
        let mut s = String::from(&data);
        let v = genvec(&mut s);
    });
}
```

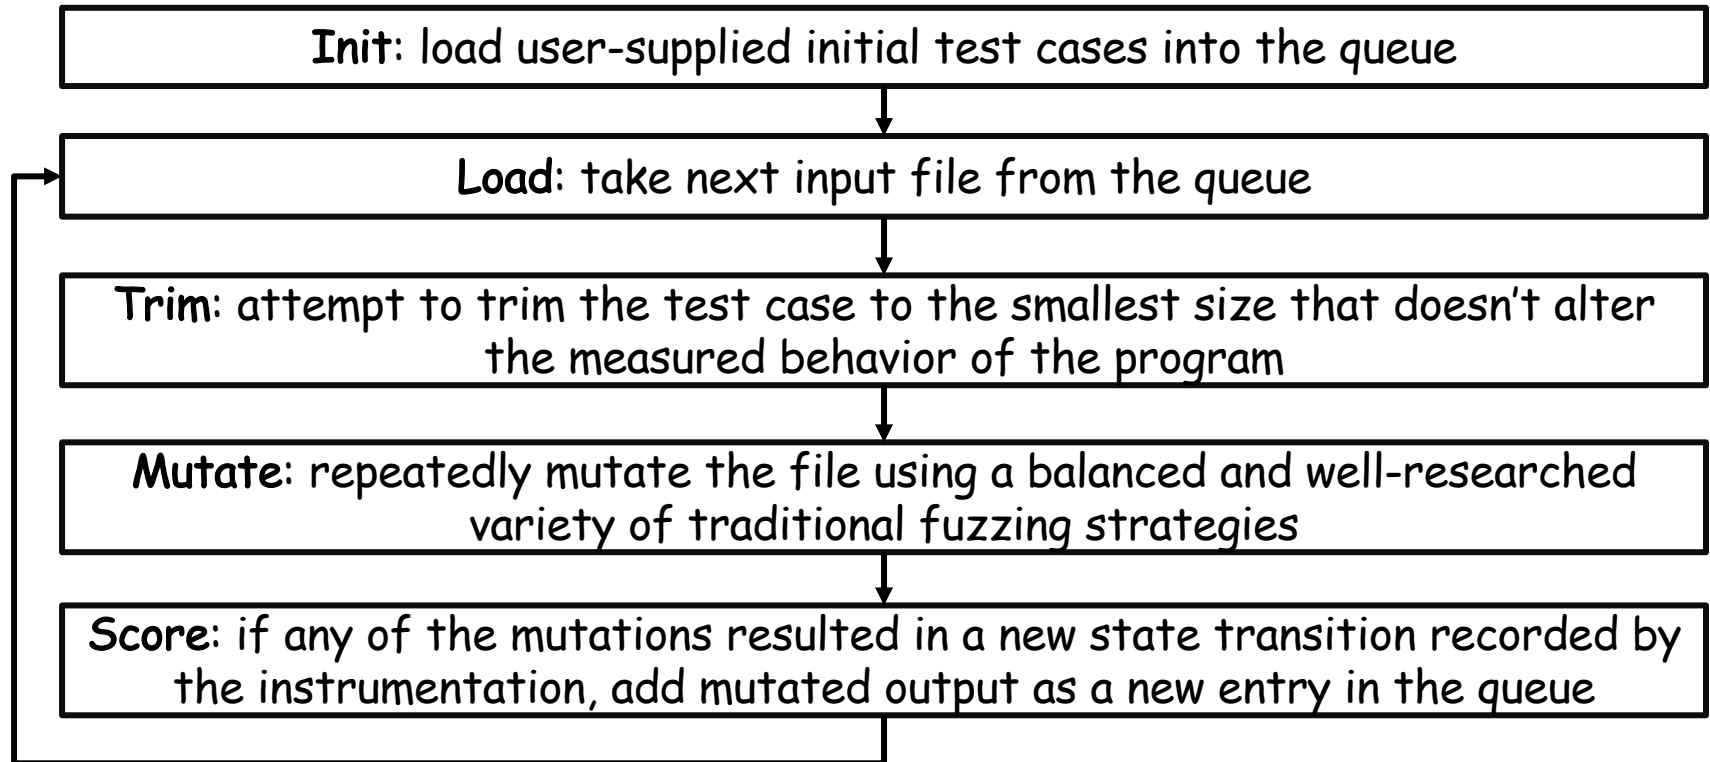
5. build the fuzz target

```
#: cargo afl build
```

6. add initial test seeds into a folder and then fuzz

```
#: cargo afl fuzz -i initcases -o output targetbin
```

Overall Algorithm



The discovered test cases are also periodically culled to eliminate ones that have been obsoleted by newer, higher-coverage finds

Instrumentation

- Supports programs
 - written in C/C++/Objective C
 - variants for Python/Go/Rust/OCaml
- Code instrumented to observe execution paths:
 - if source code is available, using modified compiler
 - if source code is unavailable, running code in an emulator
- Code coverage represented as a 64KB bitmap
 - different executions may collide with a small chance.
 - tradeoff between collision and efficiency

How to Instrument the Source Code?


- Objective: Different jumps from src to dest should result in different offsets of the bitmap.
- Where to inject?
 - at every branch point
- What information to inject?
 - insert a random number as the id/hash of the block.

Pseudocode:

```
cur_location = <COMPILE_TIME_RANDOM_FOR_THE_BLOCK>;  
bitmap[cur_location  $\oplus$  (prev_location >> 1)]++;
```


More about the Design

- Why shift 1 bit? Consider the following cases



How to differentiate

- Block A → Block B
- Block B → Block A



How to differentiate


- Block A → Block A
- Block B → Block B

- Edge coverage in nature
- Collision rate?

Branch cnt	Colliding tuples	Example targets
1,000	0.75%	giflib, lzo
2,000	1.5%	zlib, tar, xz
5,000	3.5%	libpng, libwebp
10,000	7%	libxml
20,000	14%	sqlite
50,000	30%	-

How to Instrument the Binaries


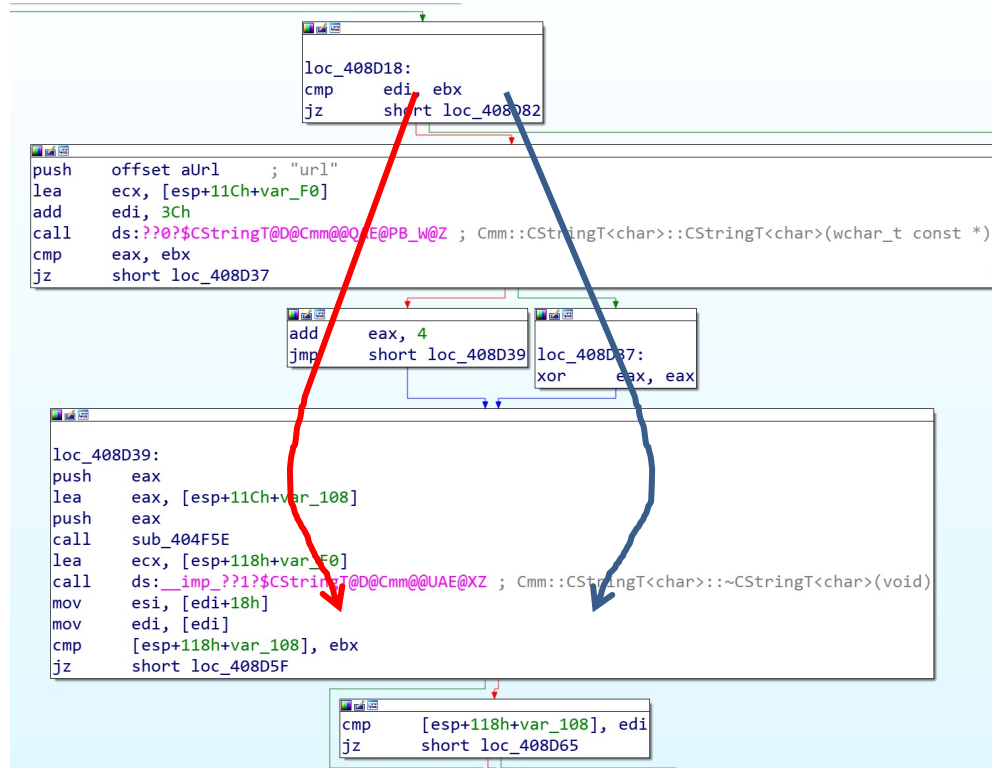
- Use the instruction address?
 - 32/64 bit
 - The bitmap offset is only 16bit;
- Select important bits heuristically.



Pseudocode:

```
cur_location = (block_addr >> 4)  $\oplus$  (block_addr << 8);  
bitmap[cur_location  $\oplus$  (prev_location >> 1)]++;
```


Visualize Bitmap




Bitmap (64KB)

Sample Mutation Strategies

- Bit flip
- Increase/decrease integers
- Use predefined interesting values
 - *eg.*, 0, -1, MAX_INT for integer
- Delete/combine/zero input block

Example: Learning the JPG file format

- Task: to fuzz a program that expects a JPG as input
- Experiment: start with 'hello world' as initial test input
- Result:
 - Generate the first image after six hours on an 8-core system
 - Produced many interesting pics by using the image as a seed



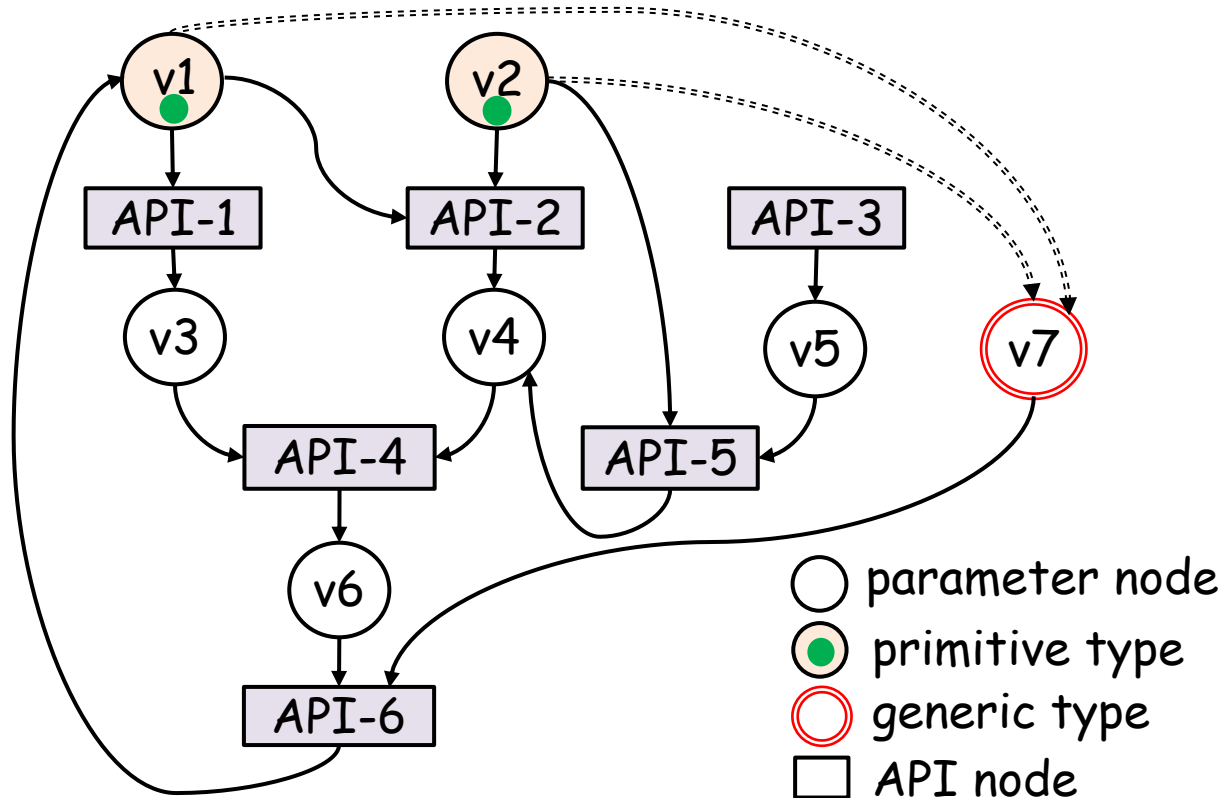
3. Fuzz Target Generation

Problem

- How to fuzz Rust libraries automatically?
- We need fuzz targets for library APIs
- How to generate fuzz target automatically?

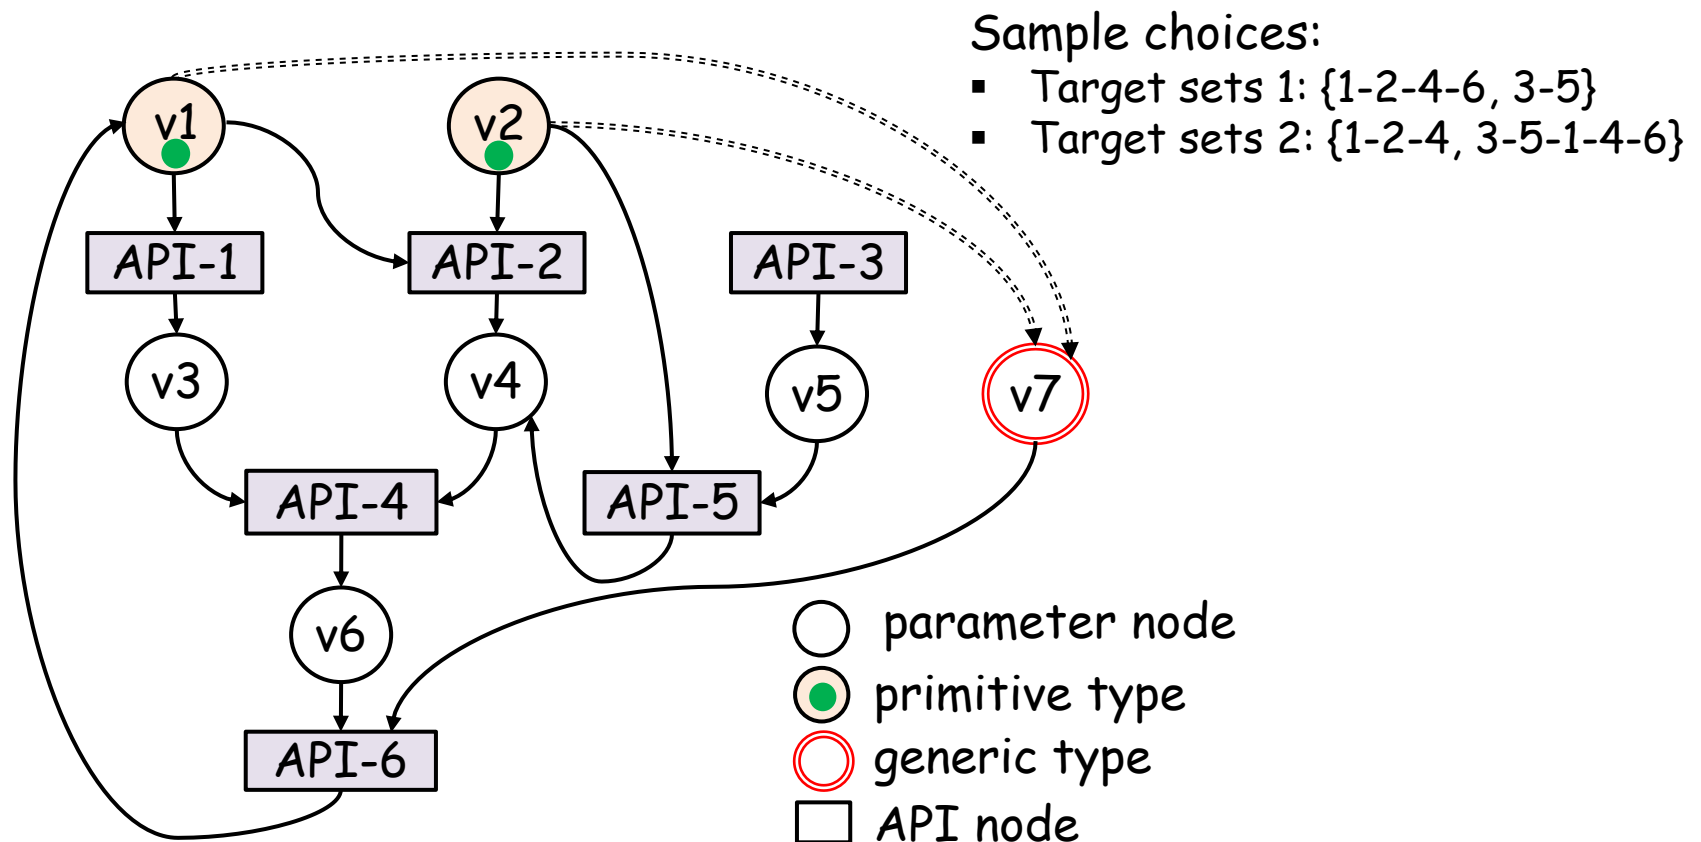
Modeling the Dependencies of APIs

- Use PetriNet



How to Traverse the Graph?

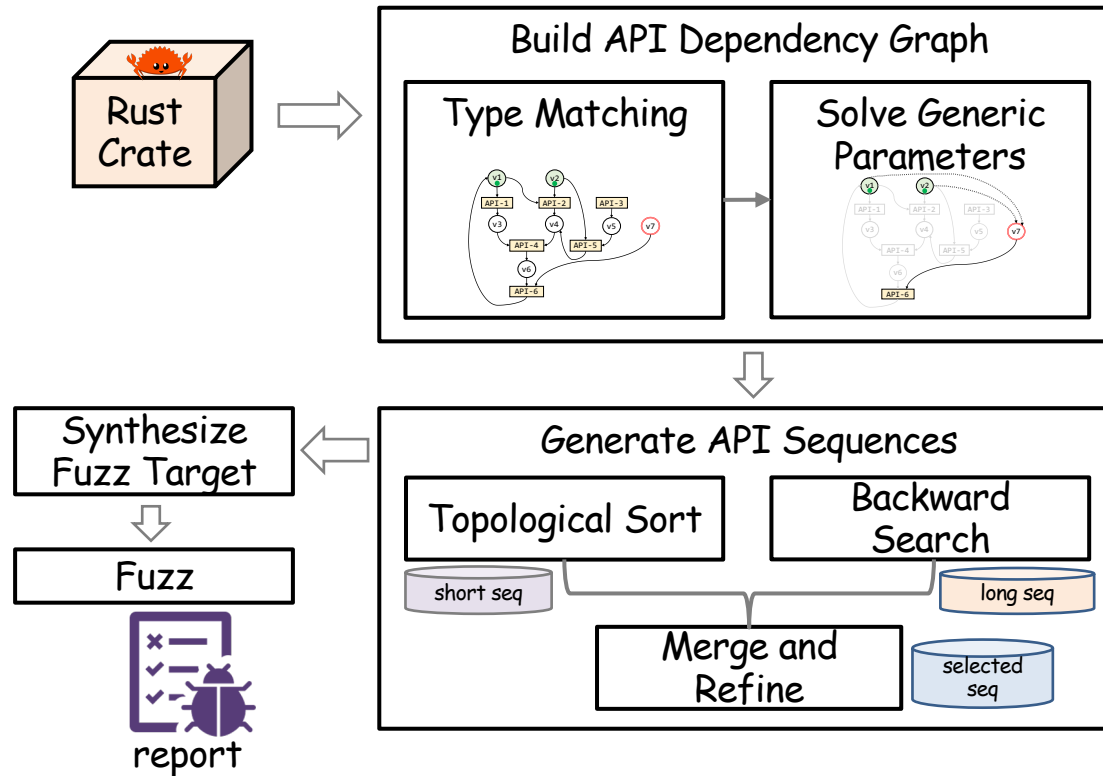
- Objective: high API coverage + efficiency
- Topological sort: generate short sequences
- Backward search: generate long sequences



Sequence Selection

- Selecting the sequence that contains the most uncovered API nodes
- Selecting the sequence that can activate the most uncovered parameter nodes
- Selecting the shortest sequence for those contribute equally using the above two rules
- Randomly picking one sequence if there are multiple candidates.

Framework Overview



4. Dynamic Symbolic Execution

Concolic (Concrete + Symbolic) Execution

Dynamic Symbolic Execution

- Stores program state concretely and symbolically
- Step:
 - 1) Start with random input values
 - 2) Keep track of both concrete values and symbolic constraints
 - 3) Use concrete values to simplify symbolic constraints
 - 4) Solve the constraints

An Illustrative Example

```
fn foo(x:usize)->usize{
    2*x;
}

fn test(x:usize, y:usize){
    let z = foo(y);
    if (z == x){
        if (x > y+10){
            unreachable!();
        }
    }
}
```

Concrete
Execution

Symbolic
Execution

concrete
state

symbolic
state

path
condition

$x = 22$

$x = x_0$

$2*y_0 \neq x_0$

$y = 7$

$y = y_0$

$z = 14$

$z = 2*y_0$

Solve: $2*y_0 == x_0$

Solution: $x_0 = 2, y_0 = 1$

An Illustrative Example

```
fn foo(x:usize)->usize{
    2*x;
}

fn test(x:usize, y:usize){
    let z = foo(y);
    if (z == x){
        if (x > y+10){
            unreachable!();
        }
    }
}
```

Concrete
Execution

Symbolic
Execution

concrete
state

symbolic
state

path
condition

$x = 2$

$x = x_0$

$2 * y_0 == x_0$

$y = 1$

$y = y_0$

$x_0 \leq y_0 + 10$

$z = 2$

$z = 2 * y_0$

Solve: $(2 * y_0 == x_0)$ and $(x_0 > y_0 + 10)$
Solution: $x_0 = 30, y_0 = 15$

An Illustrative Example

```
fn foo(x:usize)->usize{
    2*x;
}

fn test(x:usize, y:usize){
    let z = foo(y);
    if (z == x){
        if (x > y+10){
            unreachable!();
        }
    }
}
```

Concrete
Execution

Symbolic
Execution

concrete
state

symbolic
state

path
condition

x = 30

$x = x_0$

$2 * y_0 == x_0$

y = 15

$y = y_0$

$x_0 > y_0 + 10$

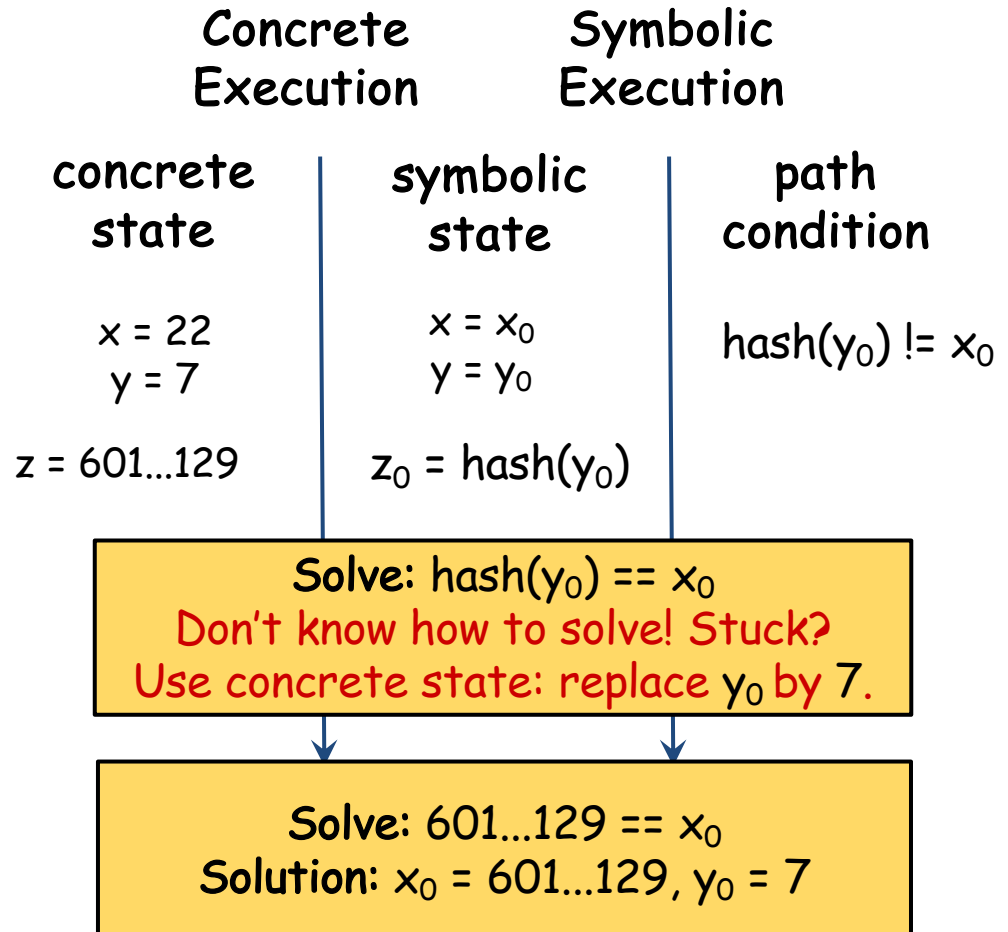
z = 30

$z = 2 * y_0$

A More Complex Example

```
fn foo(x:usize)->usize{
    hash(x);
}

fn test(x:usize, y:usize){
    → let z = foo(y);
    → if (z == x){
        if (x > y+10){
            unreachable!();
        }
    }
}
```



A Third Example

```
fn foo(x:usize)->usize{
    hash(x);
}

fn test(x:usize, y:usize){
  → let z = foo(y);
  → if (z == x){
      if (x > y+10){
          unreachable!();
      }
  }
}
```

Concrete Execution	Symbolic Execution	
concrete state	symbolic state	path condition
$x = 22$ $y = 7$	$x = x_0$ $y = y_0$	$x_0 \neq y_0$ $\text{hash}(x_0) \neq \text{hash}(y_0)$
Solve: $x_0 \neq y_0$ and $\text{hash}(x_0) == \text{hash}(y_0)$ Use concrete state: replace y_0 by 7.		
Solve: $x_0 \neq 7$ and $\text{hash}(x_0) == 601\dots129$ Use concrete state: replace x_0 by 22.		
Solve: $22 \neq 7$ and $438\dots861 == 601\dots129$ Unsatisfiable!		

SE Tools Recommended

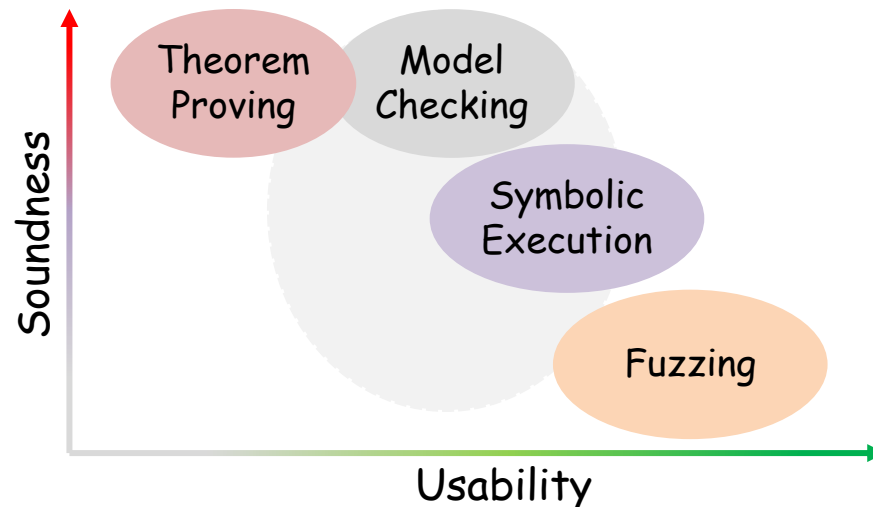
- angr: based on binary
 - <https://angr.io>
- KLEE: require source-code
 - <https://klee.github.io>

Limitation of Symbolic Execution

- Symbolic reasoning (constraint modeling/solving)
 - array/dynamic memory
 - floating-point arithmetic
 - runtime environment
 - ...
- Path-explosion (scalability): large computation tree
 - loops
 - external functions
 - ...

Comparison

- Formal methods should be sound
 - Theorem proving: prove some properties based on axioms
 - Model checking: check properties by enumerating each possible execution
 - e.g., via symbolic execution
- Fuzzing is the most easy-to-use approach



More Reference

- AFL: <https://lcamtuf.coredump.cx/afl>
- angr: <https://angr.io>
- KLEE: <https://klee.github.io>
- "RULF: Rust library fuzzing via API dependency graph traversal", ASE, 2021.
- "Benchmarking the capability of symbolic execution tools with logic bombs", TDSC, 2020.